

# Efficient, large-scale simulations on complex geometries using a hybrid multigrid solver and a fully distributed triangulation

Peter Munch<sup>1,2</sup> and Martin Kronbichler<sup>2</sup>

<sup>1</sup>Institute of Materials Research, Materials Mechanics, Helmholtz-Zentrum Geesthacht, Germany

<sup>2</sup>Institute for Computational Mechanics, Technical University of Munich, Germany

August 6, 2019

# Peter Munch (@peterrum)

- Ph.D. candidate at TUM and employee at HZG
- Supervised by Martin Kronbichler
- Research interests:
  - efficient and hardware-aware algorithms (incl. cache analysis)
  - **matrix free**, shared memory, **large-scale simulations with MPI**
  - **iterative solvers (hybrid multigrid)**
  - applications: **fluid mechanics**, solid mechanics, material science, and plasma physics
- For this workshop:
  - discuss the newly developed fully distributed triangulation
  - provide help regarding parallelization, vectorization, matrix free, ...



<https://www.lnm.mw.tum.de/staff/peter-muench/>

Part 1:

## **Problem definition and motivation**

## Definition (1)


*A geometry is complex if it can be meshed only with a significant amount of cells.*

### Definition (2)

*A geometry is complex if it can be meshed only with a significant amount of coarse cells<sup>1</sup>.*


---

<sup>1</sup>**Note:** Meshes generated by black-box mesh-generation software might also have a large amount of coarse cells. The algorithms presented here can be applied also for those meshes.




- ▶ mesh model of the first 11 airway generations of a patient-specific lung geometry
- ▶ 84,864 coarse cells + single global refinement
- ▶ goal: to understand the ventilation in the entire lung, i.e., oxygen transport and mixture between fresh and consumed air


**7 generations**



**8 generations**



**9 generations**



**coarse-grid cells:**


**4,236**

**9,396**


**16,032**

**7 generations with**


**(coarse grid)**



**1 refinement**



**2 refinements**



**cells:**

4,236

33,888

271,104

**using transfinite interpolation**



Simulations with up to **16 lung generations with at least a million coarse cells** in total are needed to be able to retrieve biologically meaningful results.

Solving the Poisson equation<sup>2</sup>:

- ▶ with  $\text{FE}_Q(k)$  and  $\text{FE}_{\text{DGQ}}(k)$  with  $k > 1$
- ▶ huge coarse grid ( $\gg 10,000$  coarse-grid cells)
- ▶ small number of refinements ( $\leq 2$  refinements)
- ▶ no adaptivity (no hanging nodes)

requires:

1. an efficient matrix-free hybrid multigrid solver and
2. a “fully distributed” triangulation.

▷ Part 2/3

▷ Part 4

---

<sup>2</sup>Solving the Poisson equation is a substep of the dual-splitting projection scheme for solving the Navier–Stokes equations and of the solution of the Vlasov-Poisson equation (plasma physics).

Part 2:

## Matrix-free geometric multigrid methods (recap)<sup>3</sup>

---

<sup>3</sup> see also: Martin Kronbichler and Wolfgang A Wall. A Performance Comparison of Continuous and Discontinuous Galerkin Methods with Fast Multigrid Solvers. SIAM Journal on Scientific Computing, 40(5):A3423A3448, jan 2018. ISSN 1064-8275. doi: 10.1137/16M110455X.

Solve the system of linear equations  $\mathcal{A}(\mathbf{x}) = \mathbf{b}$ :

- ▶ pre-smoothing:

$$\mathbf{x} \leftarrow \mathcal{S}(\mathbf{x})$$

- ▶ recursive coarse-grid correction:

$$\mathcal{A}_c(\mathbf{v}) = \mathcal{R}(\mathbf{b} - \mathcal{A}(\mathbf{x})) \quad \text{and} \quad \mathbf{x} \leftarrow \mathbf{x} + \mathcal{P}(\mathbf{v})$$



- ▶ post-smoothing:

$$\mathbf{x} \leftarrow \mathcal{S}(\mathbf{x})$$

Geometry:



↓ Fine grid



↑ Coarse grid

**needed: smoother ( $\mathcal{S}$ ), intergrid transfer operator ( $\mathcal{I}$ ,  $\mathcal{P}$ ,  $\mathcal{R}$ ), coarse-grid solver**

Chebyshev smoother with point Jacobi inner preconditioner:

$$\mathbf{u}_{j+1} = \mathbf{u}_j + \sigma_j(\mathbf{u}_j - \mathbf{u}_{j-1}) + \theta_j \mathbf{D}^{-1}(\mathbf{b} - \mathbf{A}\mathbf{u}_j)$$

with the main building block:

Matrix-free matrix-vector product

$$\mathbf{v} = \mathbf{A}\mathbf{u} = \left( \sum_c \Pi_c^T \mathbf{A}_c \Pi_c \right) \mathbf{u} = \sum_c \Pi_c^T \mathbf{A}_c \Pi_c \mathbf{u} = \boxed{\sum_c \Pi_c^T \circ \mathcal{A}_c \circ \Pi_c(\mathbf{u}) = \mathbf{v}}$$

- ▶ reduction of data to be loaded (+mapping)
- ▶ reduction of work, using sum factorization within  $\mathcal{S}$  in  $\mathcal{A}_c = \mathcal{S}^T \circ \mathcal{Q}_c \circ \mathcal{S}$
- ▶ vectorization over elements (enabled by  $\text{AoS} \leftrightarrow \text{SoA}$ -functionality of  $\Pi_c$ )
- ▶ improved cache usage

### Example: cell contribution of Laplace operator

Listing 1:  $v = \sum_c \Pi_c^T \circ \mathcal{S}^T \circ \mathcal{Q}_c \circ \mathcal{S} \circ \Pi_c(u)$


```
data.cell_loop(this, cell_loop, dst, src);
```

Listing 2:  $\Pi_c^T \circ \mathcal{S}^T \circ \mathcal{Q}_c \circ \mathcal{S} \circ \Pi_c(u)$

```
void cell_loop(data, dst, src, range) const {  
    FEEvalCell fe_eval(data);  
  
    for(auto cell = range.first; cell < range.second; ++cell) {  
        fe_eval.reinit(cell);  
        fe_eval.gather_evaluate(src, false, true);  
        this->do_cell_integral(fe_eval);  
        fe_eval.integrate_scatter(false, true, dst);  
    }  
}
```

Listing 3:  $\mathcal{Q}_c(u_q) \rightarrow$  “kernel”

```
void do_cell_integral(fe_eval) const {  
    for(unsigned int q = 0; q < fe_eval.n_q_points; ++q)  
        fe_eval.submit_gradient(fe_eval.get_gradient(q), q);  
}
```



## Example: cell contribution of Laplace operator

```
namespace MF::Util {
    class Kernel{
        Kernel(/*Flags*/) {}
        virtual void do_cell_integral(fe_eval) const {};
        virtual void do_face_integral(fe_eval_m, fe_eval_p) const {}; //DG
        virtual void do_boundary_integral(fe_eval) const {}; //DG
    }

    template<typename Kernel>
    class LinearOperator {
        LinearOperator(Kernel kernel) : kernel(kernel) {}
        void vmult(dst, src);

        Kernel & kernel;
    }
}
```

```
class LaplaceKernel : public MF::Util::Kernel {
    LaplaceKernel : MF::Util::Kernel(/*flags*/) {}


    void do_cell_integral(fe_eval) const {
        for(unsigned int q = 0; q < fe_eval.n_q_points; ++q)
            fe_eval.submit_gradient(fe_eval.get_gradient(q), q);
    }
}
```

**New infrastructure → PR????:**

- ▶ DG/CG
- ▶ multigrid levels
- ▶ ECL/FCL
- ▶ scalar/vectorial
- ▶ constraints

Due to quad-/oct-tree refinement, the global prolongation/restriction operation can be reformulated as a local prolongation/restriction operation of each coarse cell  $c$  onto its 4/8 children  $f$ <sup>4</sup>:

$$\mathbf{x}_f = \mathcal{P}_h(\mathbf{x}_c) = \sum_{(c,f)} \Pi_f^T \circ \mathcal{P}_h^{(e)} \circ \Pi_c(\mathbf{x}_c) \quad \text{and} \quad \mathbf{x}_c = \mathcal{R}_h(\mathbf{x}_f) = \mathcal{P}_h^T(\mathbf{x}_f)$$




▷ matrix-free via sum factorization, for Q and DG\_Q!

<sup>4</sup>Note: in the case of Q, degrees of freedom shared by multiple cells have to be weighted by  $\Pi$ .



**Back to the example:** 7 generations with `FE_DGQ(5)` :



Number of DoFs on coarse grid:

$$\approx 4,236 \times (5 + 1)^3 = 914,976$$

This coarse problem is not suited for:

- ▶ plain conjugate gradient methods (CG)
- ▶ algebraic multigrid (AMG)

**Remedy: hybrid multigrid solver!**

## Definition (3)

*A geometry is complex if the computation time is dominated by solving the coarse-grid problem (i.e.  $t_c \gg t_f$ ) when only  $h$ -coarsening is exploited.*

Part 3:

## Hybrid multigrid solver<sup>5</sup>

---

<sup>5</sup>developed in my Master's thesis "An efficient hybrid multigrid solver for high-order discontinuous Galerkin methods" 2018, supervised by Martin Kronbichler and Niklas Fehn (online available:

<https://mediatum.ub.tum.de/node?id=1514962>)

### Definition

*A hybrid multigrid solver exploits the structure of the problem as much as possible by both geometrical and polynomial coarsening to decrease the size of the coarse-grid problem. The resulting minimal coarse-grid problem can be solved by AMG in a black-box fashion.*

### Observation

*In the case of discontinuous Galerkin methods (DG), the problem size can be decreased even more by switching to continuous basis.*

Each multigrid level can be identified by the following triple:

$$(\text{level}, \text{degree}, \text{FE}) \in \mathbb{N}_0 \times \mathbb{N} \times \{\mathcal{Q}, \text{DGQ}\}$$

with following transition functions:

$$\begin{aligned}(\text{level}, \text{degree}, \text{FE}) &\xrightarrow{h} (\text{level}-1, \text{degree}, \text{FE}) \\(\text{level}, \text{degree}, \text{FE}) &\xrightarrow{p} (\text{level}, \mathbf{next}(\text{degree}), \text{FE}) \\(\text{level}, \text{degree}, \text{DGQ}) &\xrightarrow{c} (\text{level}, \text{degree}, \mathcal{Q}) \\(0, 1, \mathcal{Q}) &\rightarrow \perp\end{aligned}$$

Each multigrid level can be identified by the following triple:

$$(\text{level}, \text{degree}, \text{FE}) \in \mathbb{N}_0 \times \mathbb{N} \times \{\mathcal{Q}, \text{DGQ}\}$$

with the **p-coarsening strategies next()**:

$$\textcircled{1} \ k_{i-1} = 1 \qquad \textcircled{2} \ k_{i-1} = \max(1, \lfloor k_i/2 \rfloor) \qquad \textcircled{3} \ k_{i-1} = \max(1, k_i - 1)$$

Observation: bisection (strategy  $\textcircled{2}$ ) is the most efficient; it is cheap and its iteration numbers are only weakly dependent on degree.

Each multigrid level can be identified by the following triple:

$$(\text{level}, \text{degree}, \text{FE}) \in \mathbb{N}_0 \times \mathbb{N} \times \{Q, \text{DGQ}\}$$

### Example:

GMG:

$$(2, 5, \text{DGQ}) \xrightarrow{h} (1, 5, \text{DGQ}) \xrightarrow{h} (0, 5, \text{DGQ})$$

HMG:

$$(2, 5, \text{DGQ}) \xrightarrow{c} (2, 5, Q) \xrightarrow{p} (2, 2, Q) \xrightarrow{p} (2, 1, Q) \xrightarrow{h} (1, 1, Q) \xrightarrow{h} (0, 1, Q)$$

with:

- + 3 additional explicit multigrid levels
- + reduction of the size of the coarse problem by  $\approx 6^3 = 216$  for 3D
- + linear continuous elements suited very well for AMG (e.g. ML and MueLu).

Based on the observation that<sup>6</sup>:

$$(\text{level}, \underbrace{\text{degree, FE}}_{\text{def. of DoFHandler}}) \in \mathbb{N}_0 \times \mathbb{N} \times \{Q, DGQ\}$$

the following deal.II-specific implementation can be derived:


- ▶ construct a DoFHandler for each (degree, FE)-pair and
- ▶ define new multigrid transfer operators  $\mathcal{P}_\square$ ,  $\mathcal{R}_\square$ , and  $\mathcal{I}_\square$  with  $\square \in \{h, p, c\}$  working on Vectors living on different DoFHandlers.

---


<sup>6</sup>`DoFHandler<dim> dof_handler(tria); dof_handler.distribute(FE_Q(degree));`



▷ **p-transfer** (for  $\mathcal{Q}$  and DGQ; with sum factorization):




▷ **c-transfer** (for arbitrary polynomial degrees):



$$\mathbf{x}_f = \mathcal{P}_{\square}(\mathbf{x}_c) = \sum_{(c,f)} \Pi_f^T \circ \mathcal{P}_{\square}^{(e)} \circ \Pi_c(\mathbf{x}_c)$$

## Virtualization:




Requirement: AMG needs a matrix representation of the operator.

Is there a way to create a matrix  
with a matrix-free framework?

The  $i$ -th column of the local element matrix can be determined via matrix-free operator evaluations:

$$\text{column}_i(A_K) = \mathcal{A}_K(e_i)$$



A sequence of calculation of single columns leads to the full (element) matrix.

Remarks:

- ▶ matrix-free complexity:  $d \cdot (k+1)^{2 \cdot d + 1}$  vs. naïve:  $d \cdot (k+1)^{3 \cdot d}$
- ▶ perfect vectorization

```
namespace MF::Util {  
    // Kernel and LinearOperator ...  
  
    template<typename Kernel> class SparseMatrixOperator {  
        void init_system_matrix(system_matrix) const;  
        void calculate_system_matrix(system_matrix) const;  
    }  
  
    template<typename Kernel> class DiagonalOperator {  
        void calculate_diagonal(diagonal) const;  
    }  
  
    template<typename Kernel> class BlockDiagonalOperator {  
        void calculate_block_diagonal_matrices(do_invert) const;  
        void apply_block_diagonal(dst, src) const;  
    }  
  
    template<typename Kernel> class NonLinearOperator {  
        void set_solution_linearization(vector) const;  
        void evaluate_residual(dst, src) const;  
    }  
  
    // ... and many more  
}
```

## New infrastructure → PR????:

- ▶ sparse matrix (!)
- ▶ diagonal
- ▶ block diagonal
- ▶ (matrix-free) block diagonal
- ▶ non-linear operator
- ▶ tested against each other

*Matrix-free hybrid multigrid methods are efficient for large coarse grids **if** the size of the coarse problem can be reduced to a minimum, i.e., enough **explicit multigrid levels** can be constructed via geometrical and polynomial coarsening.*

This conclusion is in conflict with the observation that in the case of complex geometries only moderate number of refinements ( $\leq 2$ ) and moderate polynomial degrees ( $\leq 5$ ) are feasible.

### Future work:

- ▶ create a pull request (containing the hybrid multigrid algorithm and the extended basis operator)
- ▶ extend the hybrid multigrid algorithm to h-adaptivity (with my master student) and hp-adaptivity (with @marcfehlings)

Part 4:

## **Fully distributed triangulation**

For *complex* geometries, `parallel::distributed::Triangulation` is not optimal:

- ▶ storing the coarse grid by every process is too memory consuming
- ▶ non-optimal partitioning: discontinuous partitions




Figure: NACA 0012 airfoil (with 156 coarse cells,  $P:D:T$ , by Elias Dejene)

For *complex* geometries, `parallel::distributed::Triangulation` is not optimal:

- ▶ storing the coarse grid by every process is too memory consuming
- ▶ non-optimal partitioning: discontinuous partitions




Figure: NACA 0012 airfoil (with 156 coarse cells,  $P:F:T$ , by Elias Dejene)



PR #8418 proposes the new `parallel::fullydistributed::Triangulation`, which only stores the **locally relevant cells** and allows more flexibility regarding **partitioning** (e.g. METIS, ParMETIS, Zoltan, p4est, etc.).

### Definition

Locally relevant cells are cells:

- ▶ that are locally owned by the current process or are ghost cells on a level, or
- ▶ one of their children is locally relevant.

### Further requirements:

- ✓ extract information from `dealii::Triangulation` and from `P:D:T`
- ✓ static mesh
- ✓ hanging nodes
- ✓ geometric multigrid
- ✓ periodicity (partly implemented)
- ✓ 2D/3D and **1D (also in parallel!)**
- ✓ I/O
  
- ✗ adaptive mesh refinement (AMR)

The creation of  $P:F:T$  is centered around  $P:F:ConstructionData$ .

### Definition

*The struct  $P:F:ConstructionData$  contains:*

- ▶ *a locally relevant coarse-grid triangulation (local cells and one layer of ghost cells)*
  - ▶ *vertices*
  - ▶ *cells with material and manifold IDs*
  - ▶ *boundary IDs*
- ▶ *a mapping of the locally relevant coarse grid into the global coarse-grid triangulation*
- ▶ *information about cell refinement and cell owner (subdomain ID)*

The creation of `P:F:T` is centered around `P:F:ConstructionData`.

Two-step process:

- ▶ create `P:F:ConstructionData`
- ▶ create `P:F:Triangulation` using it

```
#include <deal.II/distributed/tria.h>

using namespace parallel::fullydistributed;

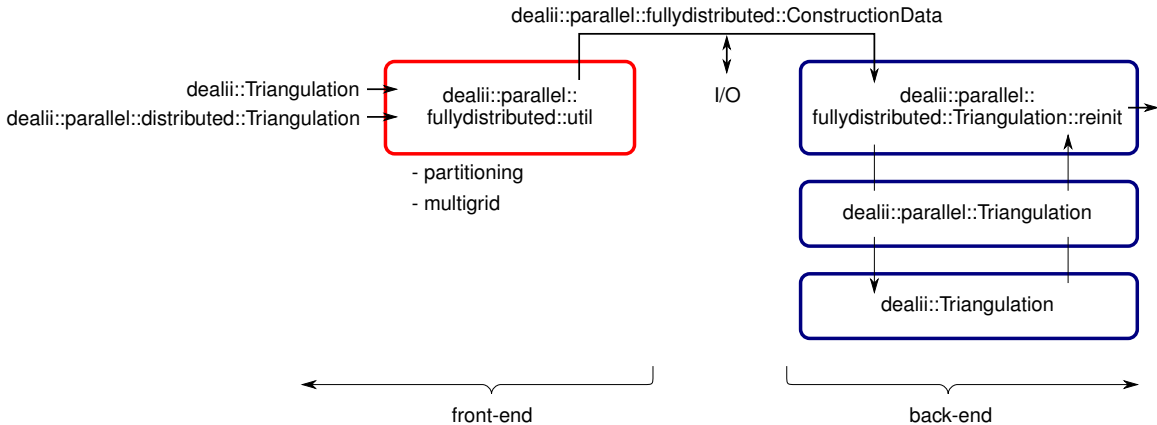
// create a construction data
ConstructionData<dim> data;
// ... and fill it (TODO: by user)

// create triangulation
Triangulation<dim> tria(comm);
tria.reinit(data); // by using the construction data
```


Creating `P:F:ConstructionData` with `P:F:Utilities`:

```
ConstructionData<dim, spacedim> copy_from_triangulation(const Triangulation<dim, spacedim> &tria);
ConstructionData<dim, spacedim> create_and_partition(std::function<void(Triangulation<dim, spacedim> &)> func);
ConstructionData<dim, spacedim> deserialize(std::string file_name);
```

The creation of  $P : F : T$  is centered around `P : F : ConstructionData`.




You can find a set of tutorials on <https://github.com/peterrum/dealii-pft>:




### Step 1: serial triangulation

▷ closes PR#3956 ping @tjhei




← colored serial mesh


↓ partitioned mesh




rank 0




rank 1



rank 2



rank 3




rank 4


### Step 2: serialization/deserialization → boost

### Step 3/4: multigrid, hanging nodes, manifolds

▷ see also deal.II step-1








### Step 5: periodic faces







## Step 7/8: P : D : T




↓ fully distributed mesh




rank 0




rank 1



rank 2



rank 3



rank 4

distributed mesh



### Step 6: Multi-level partitioning: master-slave

1. create a serial triangulation by a subset of master processes (e.g. one per node):

```
auto construction_data =  
  parallel::fullydistributed::Utilities::create_and_partition<dim, spacedim>(  
    [&](dealii::Triangulation<dim, spacedim> & tria) mutable {  
      GridGenerator::subdivided_hyper_cube(tria, n_subdivisions);  
      tria.refine_global(n_refinements);  
    }, tria_pft, additional_data);
```


2. partition the mesh using e.g. METIS, taking node-locality into account:

```
void shared_partition_triangulation(dealii::Triangulation<dim> &tria,  
                                   AdditionalData               additional_data);
```

3. distribute the  $P : F : T$  mesh information

## Future work:

- ▶ **different front-ends:** e.g. `create_and_partition` also for P:D:T
- ▶ **adaptive mesh AMR:** ▶ using `ConsensusAlgorithm` (PR #8300)



- ▶ any features missing?